

Data Manipulation with Pandas

Introduction

- Pandas: Python Data Analysis Library
- Provides rich set of functions to process various types of data.
- Provides flexible data manipulation techniques as spreadsheets and relational databases.
- An open source, providing high-performance, easy-to-use data structures and data analysis tools
- Built on the top of Numpy.
- Integrates well with matplotlib library, which makes it very handy tool for analyzing the data.
- Part of the SciPy ecosystem (Scientific Computing Tools for Python)

Data structures

- Series : One-dimensional ndarray with axis labels (including time series).
- DataFrame: Two-dimensional size-mutable, potentially heterogeneous tabular data structure with labeled axes (rows and columns). Arithmetic operations align on both row and column labels.
- The primary pandas data structure

Series

- The Series is a one-dimensional array that can store various data types, including mix data types.
- The row labels in a Series are called the index.
- Any list, tuple and dictionary can be converted in to Series using 'series' method
- Like ndarrays, the length of a Series cannot be modified after definition.
- Missing data: Represented as NaN (np.nan, a float!).
- Statistical methods from ndarray have been overridden to automatically exclude missing data.

Creating a Series

```
import pandas as pd  
s = pd.Series([9.1, 7.5, 8.63], index=['Vishnu', 'Akash', 'Aditya'],  
name='CGPA')  
print(s)
```

```
Vishnu    9.10  
Akash     7.50  
Aditya    8.63  
Name: CGPA, dtype: float64
```

Some attributes

```
import pandas as pd  
s = pd.Series([9.1, 7.5, 8.63], index=['Vishnu', 'Akash', 'Aditya'],  
name='CGPA')  
print(s.dtype)  
print(s.name)  
print(s.index)
```

float64

CGPA

Index(['Vishnu', 'Akash', 'Aditya'], dtype='object')

Creating a Series from List, Tuple, dictionary

```
import pandas as pd
h = ('Ram', '15-08-2010', 48, 3.2)
s = pd.Series(h)
print(s)

d = {'Name' : 'Ram', 'DoB' : '15-08-2010', 'Height' : 48,
'Weight' : 3.2}
ds = pd.Series(d)
print(ds)

f = ['Ram', '15-08-2020', 48, 3.2]
f = pd.Series(f, index = ['Name', 'DoB', 'Height',
'Weight'])
print(f)
```

```
0      Ram
1  15-08-2010
2      48
3      3.2
dtype: object

name      Ram
DoB      15-08-2010
Height      48
Weight      3.2
dtype: object

name      Ram
DoB      15-08-2020
Height      48
Weight      3.2
dtype: object
```

Accessing data

```
import pandas as pd
```

```
s = pd.Series([9.1, 7.5, 8.63], index=['Vishnu',  
'Akash', 'Aditya'], name='CGPA')
```

```
print(s['Vishnu'])
```

```
print(s['Akash':'Aditya'])
```

```
print(s['Akash:'])
```

9.1

Akash 7.50

Aditya 8.63

Name: CGPA, dtype: float64

Akash 7.50

Aditya 8.63

Name: CGPA, dtype: float64

Creating a View

```
import pandas as pd
s = pd.Series([9.1, 7.5, 8.63], index=['Vishnu', 'Akash', 'Aditya'], name='CGPA')
t=s['Akash:']
print(t)
t['Aditya']=9.5
print(s)
```

Akash 7.50
Aditya 8.63
Name: CGPA, dtype: float64

Vishnu 9.1
Akash 7.5
Aditya 9.5
Name: CGPA, dtype: float64

Adding two series (with automatic data alignment)

```
import pandas as pd
```

```
s = pd.Series([9.1, 7.5, 8.63],  
index=['Aditya', 'Bibek', 'Satya'],  
name='CGPA')
```

```
t = pd.Series([6, 6.3], index=['Bibek',  
'Satya'], name='Height')
```

```
u=s.add(t)
```

```
print(u)
```

```
v=s.add(t, fill_value=0)
```

```
print(v)
```

Aditya	NaN
--------	-----

Bibek	13.50
-------	-------

Satya	14.93
-------	-------

dtype: float64	
----------------	--

Aditya	9.10
--------	------

Bibek	13.50
-------	-------

Satya	14.93
-------	-------

dtype: float64	
----------------	--

DataFrame

- DataFrame can be used with two dimensional size-mutable, potentially heterogeneous tabular data structure with labeled axes
- DataFrame has two different index i.e. column-index and row-index.
- Columns can have different dtypes and can be added and removed,
- The most common way to create a DataFrame is by using the dictionary of equal-length list.
- Further, all the spreadsheets and text files are read as DataFrame.

Creating a DataFrame

```
import pandas as pd
import numpy as np
df = pd.DataFrame({'Height': [5.5, 6, 6.5], 'Weight': [np.nan,
230., 275.]},
index=['Aditya', 'Bivek', 'Vishnu'])
print(df)
print(df.dtypes)
```

	Height	Weight
Aditya	5.5	NaN
Bivek	6.0	230.0
Vishnu	6.5	275.0

Height	float64
Weight	float64
dtype: object	

Other attributes

```
print(df.shape)
```

```
print(df.size)
```

```
print(df.columns)
```

```
print(df.index)
```

```
(3, 2)
```

```
6
```

```
Index(['Height', 'Weight'], dtype='object')
```

```
Index(['Aditya', 'Bivek', 'Vishnu'], dtype='object')
```

DataFrame is by using the dictionary

```
data = { 'name' : ['AA', 'IBM', 'GOOG'],  
'date'   :   ['2001-12-01', '2012-02-10',  
'2010-04-09'], 'shares' : [100, 30, 90],  
'price' : [12.3, 10.3, 32.2]}
```

```
df = pd.DataFrame(data)
```

```
print(df)
```

```
df['owner'] = 'Unknown'
```

```
print(df)
```

	name	date	shares	price
0	AA	2001-12-01	100	12.3
1	IBM	2012-02-10	30	10.3
2	GOOG	2010-04-09	90	32.2

	name	date	shares	price	owner
0	AA	2001-12-01	100	12.3	Unknown
1	IBM	2012-02-10	30	10.3	Unknown
2	GOOG	2010-04-09	90	32.2	Unknown

```

df.index = ['one', 'two',
'three']
print(df)
df = df.set_index('name',
drop=False)
print(df)

```

	name	date	shares	price	owner
one	AA	2001-12-01	100	12.3	Unknown
two	IBM	2012-02-10	30	10.3	Unknown
three	GOOG	2010-04-09	90	32.2	Unknown

	date	shares	price	owner
name				
AA	2001-12-01	100	12.3	Unknown
IBM	2012-02-10	30	10.3	Unknown
GOOG	2010-04-09	90	32.2	Unknown

Accessing data

```
print(df['shares'])  
print(df.loc['AA',:])  
print(df.loc[:, 'name'])  
print(df.loc['AA', 'shares'])
```

```
name  
AA    100  
IBM    30  
GOOG    90  
Name: shares, dtype: int64  
name      AA  
date    2001-12-01  
shares      100  
price      12.3  
owner    Unknown  
Name: AA, dtype: object  
name  
AA    AA  
IBM   IBM  
GOOG  GOOG  
Name: name, dtype: object  
100
```


Deleting any Column

```
del df['owner']
print(df)
df.drop('shares', axis =
1,inplace = True)
print(df)
df.drop(['AA',
'IBM'],axis=0,
inplace=True)
print(df)
```

	name	date	shares	price
name				
AA	AA	2001-12-01	100	12.3
IBM	IBM	2012-02-10	30	10.3
GOOG	GOOG	2010-04-09	90	32.2

	name	date	price
name			
AA	AA	2001-12-01	12.3
IBM	IBM	2012-02-10	10.3
GOOG	GOOG	2010-04-09	32.2

	name	date	price
name			
GOOG	GOOG	2010-04-09	32.2

Summing over columns and rows

```
print(df.sum())
```

```
name          AAIBMGOOG
date    2001-12-012012-02-102010-04-09
shares                220
price                54.8
owner      UnknownUnknownUnknown
dtype: object
```

```
print(df.sum(axis=1))
```

```
name
AA    112.3
IBM   40.3
GOOG  122.2
dtype: float64
```

Reading files

```
import pandas as pd
casts = pd.read_csv('cast.csv', index_col=None)
print(casts.head())
```

```
   title  year  name  type  character  n
0  Closet Monster  2015  Buffy #1  actor  Buffy 4  31.0
1  Suuri illusioni  1985  Homo $  actor  Guests  22.0
2  Battle of the Sexes  2017  $hutter  actor  Bobby Riggs Fan  10.0
3  Secret in Their Eyes  2015  $hutter  actor  2002 Dodger Fan  NaN
4  Steve Jobs  2015  $hutter  actor  1988 Opera House Patron  NaN
```

```
titles = pd.read_csv('titles.csv', index_col = None)
print(titles.tail())
```

		title	year
49995		Rebel	1970
49996		Suzanne	1996
49997		Bomba	2013
49998	Aao Jao Ghar Tumhara		1984
49999		Mrs. Munck	1995

```
a=pd.read_csv('cast.csv', usecols= ['title','year'])  
print(a.head(6))
```

```
      title  year  
0  Closet Monster  2015  
1  Suuri illusioni  1985  
2  Battle of the Sexes  2017  
3  Secret in Their Eyes  2015  
4      Steve Jobs  2015  
5  Straight Outta Compton  2015
```

Row and column selection

```
t = titles['title']  
print(t.head(3))
```

```
0          The Rising Son  
1  The Thousand Plane Raid  
2      Crucea de piatra  
Name: title, dtype: object
```

Filter Data

- Data can be filtered by providing some boolean expression in DataFrame.

```
movies90 = titles[ (titles['year']>=1990) & (titles['year']<2000) ]  
print(movies90.head(4))
```

```
      title  year  
0  The Rising Son  1990  
2  Crucea de piatra  1993  
12  Poka Makorer Ghar Bosoti  1996  
19  Maa Durga Shakti  1999
```

Sorting

- In filtering operation, the data is sorted by index i.e. by default 'sort_index' operation is used

```
macbeth = titles[ titles['title'] == 'Macbeth'].sort_values('year')  
print(macbeth.head())
```

	title	year
4226	Macbeth	1913
17166	Macbeth	1997
25847	Macbeth	1998
9322	Macbeth	2006
11722	Macbeth	2013

Null values

- 'isnull' command returns the true value if any row of has null values.

```
c = casts  
print(c['n'].isnull().head())
```

```
0    False  
1    False  
2    False  
3     True  
4     True  
Name: n, dtype: bool
```

- To display the rows with null values, the condition must be passed in the DataFrame

```
print(c[c['n'].isnull()].head(3))
```

```
      title  year  name  type  character  n
3  Secret in Their Eyes  2015  $hutter  actor  2002 Dodger Fan NaN
4         Steve Jobs  2015  $hutter  actor  1988 Opera House Patron NaN
5 Straight Outta Compton  2015  $hutter  actor  Club Patron NaN
```

- `df.isna().any()` returns a boolean value for each column.
- If there is at least one missing value in that column, the result is True.

```
print(c.isna().any())
```

```
title    False
year     False
name     False
type     False
character False
n        True
dtype: bool
```

- `df.isna().sum()` returns the number of missing values in each column.

```
print(c.isna().sum())
```

```
title      0  
year       0  
name       0  
type       0  
character  0  
n          28966  
dtype: int64
```

Handling Missing Values

- Drop missing values
- Replace missing values

Drop missing values

- We can drop a row or column with missing values using `dropna()` function. `how` parameter is used to set condition to drop.
- `how='any'` : drop if there is any missing value
- `how='all'` : drop if all values are missing
- Furthermore, using `thresh` parameter, we can set a threshold for missing values in order for a row/column to be dropped.

- `c.dropna(axis=0, inplace=True)`
- `print(c.head())`

```
          title  ...    n
0      Closet Monster  ...  31.0
1    Suuri illusioni  ...  22.0
2    Battle of the Sexes  ...  10.0
8  Lapis, Ballpen at Diploma, a True to Life Journey  ...   9.0
10    When the Man Went South  ...   8.0

[5 rows x 6 columns]
```

Replacing missing values

- `fillna()` function of Pandas conveniently handles missing values.
- Replace missing values with a scalar: `c.fillna(2)`
- `fillna()` can also be used on a particular column: `c['n'].fillna(1)`
- Using method parameter, missing values can be replaced with the values before or after them.
- `c.fillna(axis=0, method='ffill')`

String operations

- Various string operations can be performed using '.str.' option
- `h=t[t['title'].str.startswith("Maa ")]``.head(3)`
- `print(h)`

	title	year
19	Maa Durga Shakti	1999
3046	Maa Aur Mamta	1970
7470	Maa Vaibhav Laxmi	1989

Method	Description
<code>cat()</code>	Concatenate strings
<code>split()</code>	Split strings on delimiter
<code>rsplit()</code>	Split strings on delimiter working from the end of the string
<code>get()</code>	Index into each element (retrieve i-th element)
<code>join()</code>	Join strings in each element of the Series with passed separator
<code>get_dummies()</code>	Split strings on the delimiter returning DataFrame of dummy variables
<code>contains()</code>	Return boolean array if each string contains pattern/regex
<code>replace()</code>	Replace occurrences of pattern/regex/string with some other string or the return value of a callable given the occurrence
<code>repeat()</code>	Duplicate values (<code>s.str.repeat(3)</code> equivalent to <code>x * 3</code>)

<code>pad()</code>	Add whitespace to left, right, or both sides of strings
<code>center()</code>	Equivalent to <code>str.center</code>
<code>ljust()</code>	Equivalent to <code>str.ljust</code>
<code>rjust()</code>	Equivalent to <code>str.rjust</code>
<code>zfill()</code>	Equivalent to <code>str.zfill</code>
<code>wrap()</code>	Split long strings into lines with length less than a given width
<code>slice()</code>	Slice each string in the Series
<code>slice_replace()</code>	Replace slice in each string with passed value
<code>count()</code>	Count occurrences of pattern
<code>startswith()</code>	Equivalent to <code>str.startswith(pat)</code> for each element
<code>endswith()</code>	Equivalent to <code>str.endswith(pat)</code> for each element
<code>findall()</code>	Compute list of all occurrences of pattern/regex for each string

<code>findall()</code>	Compute list of all occurrences of pattern/regex for each string
<code>match()</code>	Call <code>re.match</code> on each element, returning matched groups as list
<code>extract()</code>	Call <code>re.search</code> on each element, returning DataFrame with one row for each element and one column for each regex capture group
<code>extractall()</code>	Call <code>re.findall</code> on each element, returning DataFrame with one row for each match and one column for each regex capture group
<code>len()</code>	Compute string lengths
<code>strip()</code>	Equivalent to <code>str.strip</code>
<code>rstrip()</code>	Equivalent to <code>str.rstrip</code>
<code>lstrip()</code>	Equivalent to <code>str.lstrip</code>
<code>partition()</code>	Equivalent to <code>str.partition</code>
<code>rpartition()</code>	Equivalent to <code>str.rpartition</code>


<code>casefold()</code>	Equivalent to <code>str.casefold</code>
<code>upper()</code>	Equivalent to <code>str.upper</code>
<code><u>find()</u></code>	Equivalent to <code>str.find</code>
<code>rfind()</code>	Equivalent to <code>str.rfind</code>
<code>index()</code>	Equivalent to <code>str.index</code>
<code>rindex()</code>	Equivalent to <code>str.rindex</code>
<code>capitalize()</code>	Equivalent to <code>str.capitalize</code>
<code>swapcase()</code>	Equivalent to <code>str.swapcase</code>
<code>normalize()</code>	Return Unicode normal form. Equivalent to <code>unicodedata.normalize</code>
<code>translate()</code>	Equivalent to <code>str.translate</code>
<code>isalnum()</code>	Equivalent to <code>str.isalnum</code>
<code>isalpha()</code>	Equivalent to <code>str.isalpha</code>

<code>isalpha()</code>	Equivalent to <code>str.isalpha</code>
<code>isdigit()</code>	Equivalent to <code>str.isdigit</code>
<code>isspace()</code>	Equivalent to <code>str.isspace</code>
<code>islower()</code>	Equivalent to <code>str.islower</code>
<code>isupper()</code>	Equivalent to <code>str.isupper</code>
<code>istitle()</code>	Equivalent to <code>str.istitle</code>
<code>isnumeric()</code>	Equivalent to <code>str.isnumeric</code>
<code>isdecimal()</code>	Equivalent to <code>str.isdecimal</code>

- Total number of occurrences can be counted using 'value_counts()' option.
- In following code, total number of movies are displayed base on years.
- `t['year'].value_counts().head()`


Plots

- `import matplotlib.pyplot as plt`
- `t = titles`
- `p = t['year'].value_counts()`
- `p.sort_index().plot()`
- `p.show()`



Grouping

- Groupby with column-names
- `c = casts`
- `cg = c.groupby(['year']).size()`
- `cg.plot()`
- `plt.show()`



groupby option can take multiple parameters for grouping

- `c = casts`
- `cf = c[c['name'] == 'Aaron Abrams']`
- `ct=cf.groupby(['year', 'title']).size().head()`
- `print(ct)`

```
year title
2003 The In-Laws      1
2004 Resident Evil: Apocalypse  1
      Siblings        1
2005 Cinderella Man   1
      Sabah           1
dtype: int64
```

- grouping based on maximum ratings in a year;
- `c.groupby(['year']).n.max().head()`
- To check the mean rating each year,
- `c.groupby(['year']).n.mean().head()`

Unstack

- we want to compare and plot the total number of actors and actresses in each decade.
- we need to group the data based on 'type'
- `c = casts`
- `c_decade = c.groupby(['type', c['year']//10*10]).size()`
- `print(c_decade)`


type	year
actor	1910 340
	1920 590
	1930 1364
	1940 1253
	1950 1490
	1960 1879
	1970 2191
	1980 2874
	1990 4051
	2000 6787
	2010 7259
actress	1910 267
	1920 353
	1930 511
	1940 528
	1950 665
	1960 702
	1970 1015
	1980 1686
	1990 2115
	2000 3872
	2010 4243

dtype: int64

- `us=c_decade.unstack()`
- `print(us)`

```
year  1910 1920 1930 1940 1950 1960 1970 1980 1990
2000 2010
type
actor   340  590 1364 1253 1490 1879 2191 2874 4051
6787 7259
actress 267  353  511  528  665  702 1015 1686 2115
3872 4243
```

- `us.plot(kind='bar')`
- `plt.show()`



Time series

- A series of time can be generated using 'date_range' command.
- 'periods' is the total number of samples;
- freq = 'M' represents that series must be generated based on 'Month'.
- By default, pandas consider 'M' as end of the month.
- Use 'MS' for start of the month.

- `rng = pd.date_range('2011-03-01 10:15', periods = 10, freq = 'M')`
- `print(rng)`

```
DatetimeIndex(['2011-03-31 10:15:00', '2011-04-30 10:15:00',  
              '2011-05-31 10:15:00', '2011-06-30 10:15:00',  
              '2011-07-31 10:15:00', '2011-08-31 10:15:00',  
              '2011-09-30 10:15:00', '2011-10-31 10:15:00',  
              '2011-11-30 10:15:00', '2011-12-31 10:15:00'],  
              dtype='datetime64[ns]', freq='M')
```

Support for time zone representation, converting to another time zone, and converting between time span representations.

Categoricals

- Similar to categorical variables used in statistics.
- Practical for saving memory and sorting data.
- “Examples are gender, social class, blood type, country affiliation”
- `s = pd.Series(["a","b","c","a"], dtype="category")`
- `print(s)`

```
dtype: category
Categories (3, object): [a, b, c]
```

References

- https://pandas.pydata.org/pandas-docs/stable/user_guide/
- GitHub awesome-pandas
- Pandas Guide by Meher Krishna Patel
- Manipulating and analyzing data with pandas by Céline Comte